

1) Linear Regression

2023-03-19

Pre-processing the data-set

```
data <- read.csv("Linear_Regression_Dataset_new.csv", header = TRUE)
processed_data <- na.omit(data)
head(processed_data)
```

```
##      x      y
## 1 24 21.54945
## 2 50 47.46446
## 3 15 17.21866
## 4 38 36.58640
## 5 87 87.28898
## 6 36 32.46387
```

```
summary(processed_data)
```

```
##           x           y
##  Min.      : 0.00   Min.      : -3.84
## 1st Qu.: 25.00   1st Qu.: 25.19
##  Median : 50.00   Median : 49.93
##  Mean   : 50.29   Mean     : 50.32
## 3rd Qu.: 74.50   3rd Qu.: 74.48
##  Max.    :100.00   Max.     :108.87
```

```
str(processed_data)
```

```
## 'data.frame':   999 obs. of  2 variables:
## $ x: num  24 50 15 38 87 36 12 81 25 5 ...
## $ y: num  21.5 47.5 17.2 36.6 87.3 ...
## - attr(*, "na.action")= 'omit' Named int 214
## ..- attr(*, "names")= chr "214"
```

```
nrow(processed_data)
```

```
## [1] 999
```

Splitting the model

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
indexs = createDataPartition(processed_data$y, times = 1, p = 0.7, list = F)
#times = no. of times to be split
#p = percentage of data to be used for training, here 70% is used of training and 30%
for testing

train = processed_data[indexs, ]
nrow(train)
```

```
## [1] 700
```

```
test = processed_data[-indexs, ]
nrow(test)
```

```
## [1] 299
```

Creating the model

```
cor(train$x, train$y)
```

```
## [1] 0.995206
```

```
#y = dependent
#x = independent
#y = slope * x + intercept
# dependent ~ independent
model <- lm(y ~ x, data = train)
model
```

```
##
## Call:
## lm(formula = y ~ x, data = train)
##
## Coefficients:
## (Intercept)          x
##    -0.2486      1.0058
```

```
summary(model)
```

```
##
## Call:
## lm(formula = y ~ x, data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.5202 -1.8853 -0.0699  1.8715  8.1150
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.248633   0.215237  -1.155    0.248
## x           1.005801   0.003741 268.842 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.816 on 698 degrees of freedom
## Multiple R-squared:  0.9904, Adjusted R-squared:  0.9904
## F-statistic: 7.228e+04 on 1 and 698 DF,  p-value: < 2.2e-16
```

Predicting the values using the model

```
#df <- data.frame(x = c(29)), just to initially check
predicted <- predict(model, test)
predicted
```

##	5	10	13	16	17	18
##	87.2560222	4.7803703	23.8905824	60.0994051	25.9021836	73.1748134
##	21	22	23	25	35	45
##	68.1458102	87.2560222	58.0878039	84.2386203	60.0994051	35.9601900
##	49	50	54	61	62	68
##	18.8615792	32.9427881	59.0936045	58.0878039	74.1806140	-0.2486328
##	69	72	73	74	75	79
##	11.8209748	64.1226077	4.7803703	58.0878039	31.9369874	4.7803703
##	82	90	93	96	115	116
##	2.7687691	95.3024273	41.9949938	55.0704020	40.9891931	26.9079843
##	117	120	121	126	127	129
##	58.0878039	70.1574115	71.1632121	45.0123957	40.9891931	36.9659906
##	132	136	141	142	146	147
##	47.0239969	95.3024273	56.0762026	80.2154178	95.3024273	94.2966267
##	149	152	157	158	166	168
##	6.7919716	48.0297976	25.9021836	51.0471995	73.1748134	80.2154178
##	169	174	175	177	183	189
##	77.1980159	56.0762026	47.0239969	1.7629684	40.9891931	16.8499779
##	196	200	205	207	209	212
##	71.1632121	51.0471995	2.7687691	23.8905824	58.0878039	13.8325760
##	217	222	228	232	236	241
##	71.1632121	74.1806140	71.1632121	99.3256298	65.1284083	28.9195855
##	243	246	248	249	255	256
##	6.7919716	50.0413988	89.2676235	76.1922153	66.1342089	50.0413988
##	258	259	270	277	283	286
##	60.0994051	34.9543893	100.3314305	37.9717912	44.0065950	4.7803703
##	293	294	301	302	303	305
##	55.0704020	3.7745697	75.1864146	33.9485887	37.9717912	20.8731805
##	306	309	310	312	313	320
##	88.2618229	44.0065950	9.8093735	15.8441773	31.9369874	34.9543893
##	321	322	326	327	329	332
##	23.8905824	16.8499779	89.2676235	69.1516108	5.7861710	8.8035729
##	334	335	343	346	349	354
##	50.0413988	85.2444210	18.8615792	58.0878039	19.8673798	98.3198292
##	355	357	358	360	363	367
##	90.2734241	93.2908260	21.8789811	12.8267754	30.9311868	22.8847817
##	369	372	374	383	387	390
##	88.2618229	91.2792248	9.8093735	26.9079843	46.0181963	8.8035729
##	393	400	406	407	417	418
##	0.7571678	91.2792248	23.8905824	48.0297976	41.9949938	95.3024273
##	425	431	432	433	434	437
##	25.9021836	23.8905824	70.1574115	28.9195855	76.1922153	87.2560222
##	438	442	443	446	449	451
##	8.8035729	16.8499779	49.0355982	89.2676235	47.0239969	4.7803703
##	452	455	467	472	473	481
##	68.1458102	40.9891931	50.0413988	18.8615792	94.2966267	50.0413988
##	483	486	490	492	494	499
##	29.9253862	15.8441773	62.1110064	29.9253862	62.1110064	20.8731805
##	501	504	505	510	518	520
##	97.3140286	47.0239969	98.3198292	17.8557786	91.2792248	85.2444210
##	527	532	534	540	542	546
##	36.9659906	68.1458102	4.7803703	25.9021836	60.0994051	28.9195855
##	557	558	562	563	564	566
##	20.8731805	98.3198292	30.9311868	93.2908260	67.1400096	24.8963830
##	569	577	588	591	592	594


```
## 20.8731805 67.1400096 49.0355982 70.1574115 91.2792248 54.0646014
##          595          596          598          599          601          617
## 38.9775919 91.2792248 21.8789811  1.7629684 65.1284083 49.0355982
##          619          621          622          623          625          627
## 46.0181963 89.2676235 36.9659906 28.9195855 96.3082279 74.1806140
##          628          631          632          633          634          637
## 34.9543893 56.0762026 17.8557786 100.3314305 54.0646014 81.2212184
##          640          642          643          645          646          649
##  0.7571678 13.8325760 24.8963830 98.3198292 97.3140286 88.2618229
##          660          663          665          668          674          678
##  3.7745697 34.9543893  9.8093735 58.0878039 61.1052058 43.0007944
##          681          683          687          696          705          709
## 45.0123957 33.9485887 67.1400096 58.0878039 35.9601900 19.8673798
##          710          714          718          723          729          731
##  4.7803703 62.1110064 13.8325760 89.2676235 93.2908260 99.3256298
##          738          739          744          752          754          757
## 69.1516108 27.9137849 76.1922153 88.2618229 30.9311868 38.9775919
##          758          761          762          768          770          776
## 64.1226077 12.8267754 72.1690127 50.0413988 12.8267754 27.9137849
##          777          781          782          791          793          794
## 81.2212184 68.1458102 26.9079843 67.1400096 63.1168070 91.2792248
##          796          801          805          810          811          812
## 13.8325760  1.7629684 41.9949938 -0.2486328 40.9891931 15.8441773
##          813          815          816          817          818          820
## 94.2966267 66.1342089 23.8905824 16.8499779 90.2734241 -0.2486328
##          821          823          827          829          830          833
## 64.1226077 98.3198292 78.2038165 89.2676235 28.9195855 11.8209748
##          835          843          847          852          853          858
## 27.9137849 73.1748134 99.3256298 31.9369874 94.2966267 93.2908260
##          862          866          867          870          876          880
## 46.0181963 43.0007944 95.3024273 34.9543893 31.9369874 72.1690127
##          884          892          895          896          898          901
## 91.2792248 73.1748134 61.1052058 99.3256298 72.1690127 78.2038165
##          912          913          915          918          920          925
## 45.0123957 59.0936045 22.8847817 95.3024273  3.7745697 96.3082279
##          927          928          929          930          933          935
## 100.3314305 87.2560222 13.8325760 13.8325760 88.2618229 65.1284083
##          938          939          943          945          946          947
## 15.8441773  4.7803703 54.0646014 44.0065950 30.9311868 68.1458102
##          949          952          957          963          964          965
## 90.2734241 95.3024273 89.2676235 45.0123957 73.1748134 57.0820032
##          966          969          971          974          977          983
## 19.8673798 55.0704020 55.0704020 72.1690127 96.3082279 64.1226077
##          988          989          991          995          999
## 41.9949938 43.0007944 92.2850254  7.7977722 62.1110064
```

Plotting the linear regression curve

```
library(ggplot2)
```


```
ggplot(processed_data, aes(x = x, y = y)) + geom_point() + geom_smooth(method = 'lm')
```

```
## `geom_smooth()` using formula = 'y ~ x'
```




```
scatter.smooth(x = processed_data$x, y = processed_data$y, main = "y ~ x")
```

y ~ x




```
plot(processed_data$x, processed_data$y, main = "y ~ x")  
abline(lm(processed_data$y ~ processed_data$x, data = processed_data), col = "blue")
```




```
res <- model$residuals  
#create Q-Q plot for residuals  
qqnorm(res)
```

Normal Q-Q Plot



```
res1 <- model$residuals  
  
plot.new()  
#text(x = 0.5, y = 0.5, labels = "This is a new plot canvas.")  
qqline(res1)
```

Confusion Matrix: *It is not used in Linear Regression as Confusion Matrix is only applied to classification problems, not regression problems.*

Conclusion: *As we can see, the model is performing poorly as the dataset is containing less number of numerical data and fewer relationships among them. As we can observe in the correlation coefficient calculation, the relationship between age and charges is not more than 50%. To improve the model, we can convert the categorical columns into numerical dummy data i.e. convert them into numerical factors which can be used to calculate in the regression analysis. From the summary of the model, we can observe that age and bmi attributes have a higher impact or significance on the overall accuracy of the model. A confusion matrix cannot be applied in regression analysis. Innovation in this experiment includes using regression along with caret package to find out major insights into the field of insurance charges and analysing the behaviour of insurance charges according to the given attributes.*